




Data Science Dojo

Regression

Agenda

- Introduction
 - Cost Function & Gradient Descent
 - Minimization
 - Implementation
 - Hands-on Example
 - Evaluating Regression Models
 - Regularization
- 

INTRODUCTION




Regression



Sales Forecasts



Housing Price
Predictions



Daily Temperature
Highs & Lows

Regression vs Classification

- Classification
 - Target is discrete with finite value set
 - **Examples:** survived/dead, face/non-face, fraud/non-fraud
- Regression
 - Target is continuous
 - **Examples:** price, weight, height, temperature

Notation: Titanic Dataset

| Passenger Id | Survived | Pclass | Name | Sex | Age | SibSp | Parch | Ticket | Fare | Cabin | Embarked |
|--------------|----------|--------|-----------------------------------------------------|--------|-----|-------|-------|------------------|---------|-------|----------|
| 1 | 0 | 3 | Braund, Mr. Owen Harris | male | 22 | 1 | 0 | A/5 21171 | 7.25 | | S |
| 2 | 1 | 1 | Cumings, Mrs. John Bradley (Florence Briggs Thayer) | female | 38 | 1 | 0 | PC 17599 | 71.2833 | C85 | C |
| 3 | 1 | 3 | Heikkinen, Miss. Laina | female | 26 | 0 | 0 | STON/O2. 3101282 | 7.925 | | S |
| 4 | 1 | 1 | Futrelle, Mrs. Jacques Heath (Lily May Peel) | female | 35 | 1 | 0 | 113803 | 53.1 | C123 | S |
| 5 | 0 | 3 | Allen, Mr. William Henry | male | 35 | 0 | 0 | 373450 | 8.05 | | S |

x_4^5

5: The passenger is in the 5th row

4: The passenger's name is the 4th column

Notation: Ozone Dataset

So how do we describe all the rows?

Row 1
Row 2
Row 3

| ozone | radiation | temperature | wind |
|-------|-----------|-------------|------|
| 41 | 190 | 67 | 7.4 |
| 36 | 118 | 72 | 8.0 |
| 12 | 149 | 74 | 12.6 |
| 18 | 313 | 62 | 11.5 |
| 23 | 299 | 65 | 8.6 |
| 19 | 99 | 59 | 13.8 |

$$x^1 = [190, 67, 7.4]$$

$$x^2 = [118, 72, 8.0]$$

$$x^3 = [149, 74, 12.6]$$

Notation: Ozone Dataset

The ozone dataset uses radiation, temperature and wind to predict ozone levels.

| | | x_1 | x_2 | x_3 | |
|-----|-------|-----------|-------------|-------|-----|
| | ozone | radiation | temperature | wind | |
| Y | 41 | 190 | 67 | 7.4 | X |
| | 36 | 118 | 72 | 8.0 | |
| | 12 | 149 | 74 | 12.6 | |
| | 18 | 313 | 62 | 11.5 | |
| | 23 | 299 | 65 | 8.6 | |
| | 19 | 99 | 59 | 13.8 | |

Using this notation, we can describe all the columns of the dataset.

Notation Summary


| | | |
|-------|----------------------------------|----------|
| x^i | Each row of features | Features |
| x_j | Each column of features | |
| X | Set of all the feature columns | |
| y^i | Each row of the target | Target |
| Y | The target column | |
| n | Number of rows in the dataset | |
| m | Number of columns in the dataset | |

COST FUNCTION AND GRADIENT DESCENT



What is a good regression line?

- Wind Speed=15 mph
- Ozone = ?
- Use the line that is **somewhere in the middle**
- How do we define "middle"?



$$h_{\theta}(x) = \theta_0 + \theta_1 x$$


Defining a line

How do we define a line in slope-intercept notation?

- $y = mx + b$

In θ notation

- $h_{\theta}(x) = \theta_1 x + \theta_0$




More Features

| y | x_1 | x_2 | x_3 |
|-------|-----------|-------------|-------|
| ozone | radiation | temperature | wind |
| 41 | 190 | 67 | 7.4 |
| 36 | 118 | 72 | 8.0 |
| 12 | 149 | 74 | 12.6 |
| 18 | 313 | 62 | 11.5 |
| 23 | 299 | 65 | 8.6 |
| 19 | 99 | 59 | 13.8 |

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

Residuals (or "Errors")

Difference between hypothesis $h_{\theta}(\mathbf{x})$ (predicted value) and true value (known target)




Cost Function

Minimize the 'cost' or 'loss' function – $J(\theta)$

- Smaller for lower error
- Larger for higher error


$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2$$




Cost Function

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$


$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2$$



$\theta_0=0$




Each point on the parabola corresponds to a line on the graph on the left



Cost function in three dimensions

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2$$



HOW DO WE FIND THE MINIMUM OF THE COST FUNCTION



Maximum/Minimum Problem

Find **two non-negative** numbers whose **sum is 9** and so that the product of one number and the square of the other number is a **maximum**.

$$\boxed{x + y = 9}$$

$$P = xy^2 \rightarrow \text{constraint}$$

Solution (1/2)

Sum of number is 9

$$9 = x + y$$

Product of two numbers is

$$\begin{aligned} P &= x y^2 \\ &= x (9-x)^2 \end{aligned}$$

Solution (2/2)

Using the product rule and chain rule from Calculus 101:

$$\begin{aligned} P' &= x (2) (9-x)(-1) + (1) (9-x)^2 \\ &= (9-x) [-2x + (9-x)] \\ &= (9-x) [9-3x] \\ &= (9-x) (3) [3-x] \\ &= 0 \end{aligned}$$

$$x=9 \text{ or } x=3$$

Maximum Problem

There are **50 apple trees** in an orchard.

Each tree produces **800 apples**. For each additional tree planted in the orchard, the apple output per tree drops by **10 apples**.

Question: *How many additional trees should be planted in the existing orchard in order to maximize the apple output of the orchard?*

Solution

$$A = (50 + t) \times (800 - 10t)$$

$$A = 40,000 + 300t - 10t^2$$

Solve for A' and set to 0 to find maximum.

$$A' = -20t + 300 = 0$$

$$t = 15$$

Adding 15 trees will maximize apple production

Gradient Descent

- Goal : minimize $J(\theta)$
- Start with some initial θ and then perform an update on each θ_j in turn:

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\theta^k)$$


- Repeat until θ converges

Gradient Descent

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\theta^k)$$

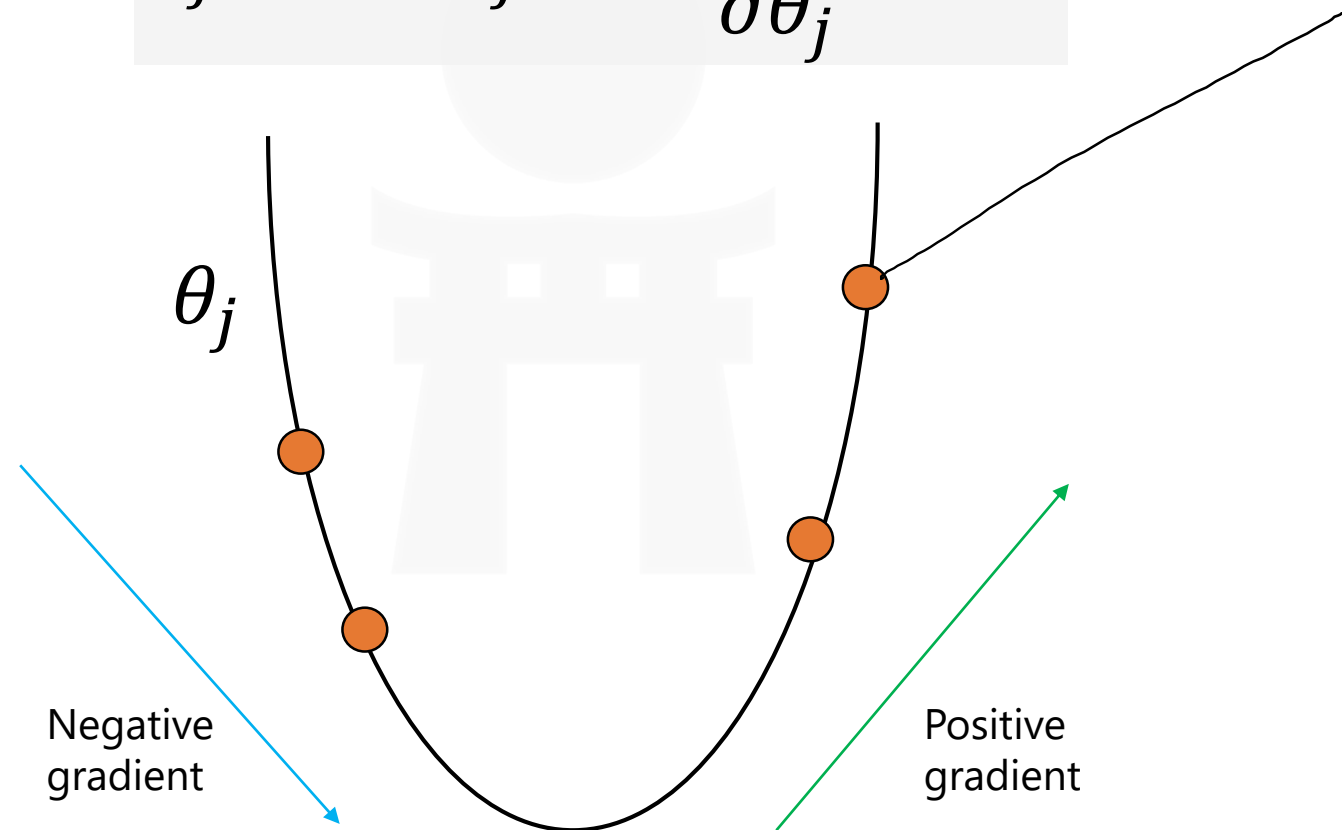
- α is known as the learning rate; set by user
- Each time the algorithm takes a step in the direction of the steepest descent and $J(\theta)$ decreases.
- α determines how quickly or slowly the algorithm will converge to a solution

Intuition



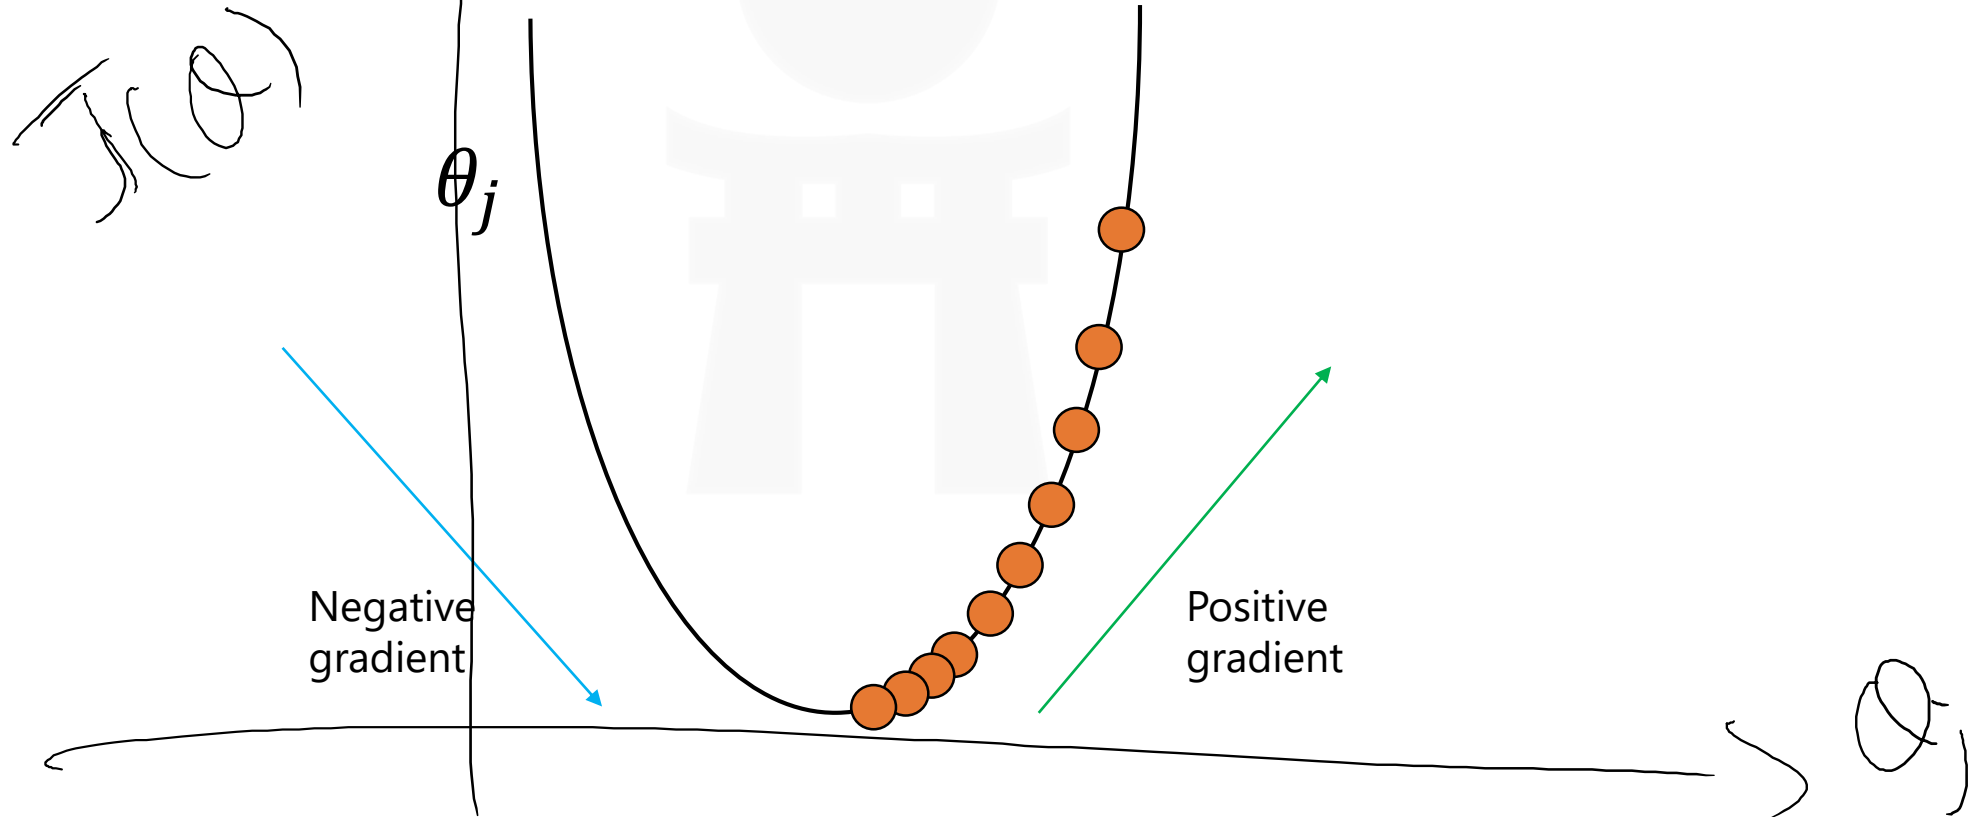
Effect of High Learning Rate: Large α

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\theta^k)$$




Learning Rate Effects Small α

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\theta^k)$$



Gradient Descent Implementation

When do we stop updating?



- When θ_j^{k+1} is close to θ_j^k
- When $J(\theta^{k+1})$ is close to $J(\theta^k)$ [Error does not change]

Batch Gradient Descent

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\theta^k)$$

Each θ_j represents one feature

- How do we incorporate all our data?
- Loop!

For j from 0 to m :

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i) x_j^i$$

- h_{θ} is updated only once the loop has completed
- Weaknesses?

Batch Gradient Descent

- Loop!

For j from 0 to m :

| ozone | radiation | temperature | wind |
|-------|-----------|-------------|------|
| 41 | 190 | 67 | 7.4 |
| 36 | 118 | 72 | 8.0 |
| 12 | 149 | 74 | 12.6 |
| 18 | 313 | 62 | 11.5 |
| 23 | 299 | 65 | 8.6 |
| 19 | 99 | 59 | 13.8 |

$$\theta_j^{k+1} := \theta_j^k - \alpha \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i) x_j^i$$

Stochastic Gradient Descent

- Consider an alternative approach:

for i from 1 to n:

for j from 0 to m:

$$\theta_j^{k+1} := \theta_j^k - \alpha(h_\theta(x^i) - y^i)x_j^i$$

- h_θ is updated when inner loop is complete
- If the training set is big, converges quicker than batch
- May oscillate around a minimum of $J(\theta)$ and never converge

* We're now only taking one random observation at a time as a sample, instead of averaging across observations

Batch vs. Stochastic

Which is the better to use? It depends.

| | Batch Gradient Descent | Stochastic Gradient Descent |
|--------------------------|----------------------------------------------|--------------------------------------------------------------|
| Function | Updates hypothesis by scanning whole dataset | Updates hypothesis by scanning one training sample at a time |
| Rate of convergence | Slowly | Quickly (but may oscillate at minimum) |
| Appropriate Dataset Size | Small | Large |

EVALUATING REGRESSION MODELS



Evaluation metrics for regression

- Mean Absolute Error (MAE)
- Root-Mean-Square Error (RMSE)
 - Root-Mean-Square Deviation
- Coefficient of Determination (R^2)

Mean Absolute Error

$$MAE(\theta) = \frac{\sum_{i=1}^n |h_{\theta}(x^i) - y^i|}{n}$$

- Mean of residual values
- "Pure" measure of error

Mean Absolute Error - Example

$$y = \{36, 19, 34, 6, 1, 45\}$$

$$h_{\theta}(x) = \{27, -2.6, 13, -7.3, -2.6, 48\}$$

$$|h_{\theta}(x) - y| = \{9, 21.6, 21, 13.3, 3.6, 3\}$$

$$MAE(\theta) = \frac{71.5}{6} = 11.9$$

Root-Mean-Square Error

$$RMSE(\theta) = \sqrt{\frac{\sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2}{n}}$$

- Square root of mean of squared residuals
- Penalizes large errors more than small
- Good measure to use to accentuate outliers

RMSE - Example

create a handout
using slides 50 — 52
to show difference b/w
RMSE & MAE
for two examples.

$$y = \{36, 19, 34, 6, 1, 45\}$$

$$h_{\theta}(x) = \{27, -2.6, 13, -7.3, -2.6, 48\}$$

$$(h_{\theta}(x) - y)^2 = \{81, 467, 441, 177, 13, 9\}$$

$$RSME(\theta) = \sqrt{\frac{1187}{6}} = 14.1$$

Coefficient of Determination (R^2)

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

where

$$SS_{res} = \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 \quad SS_{tot} = \sum_{i=1}^n (y^i - \bar{y})^2$$


SS_{res} – Sum of squared residuals (i.e. total squared error)

SS_{tot} – Sum of squared differences from mean (i.e. total variation in dataset)


Result: Measure of how well the model explains the data

- "Fraction of variation in data explained by model"

Coefficient of Determination




SS_{tot}



SS_{res}

R² Example


- $R^2 = 0.277$
- Want a much better model for real application
- $R^2 = 0.6$ can be a good model



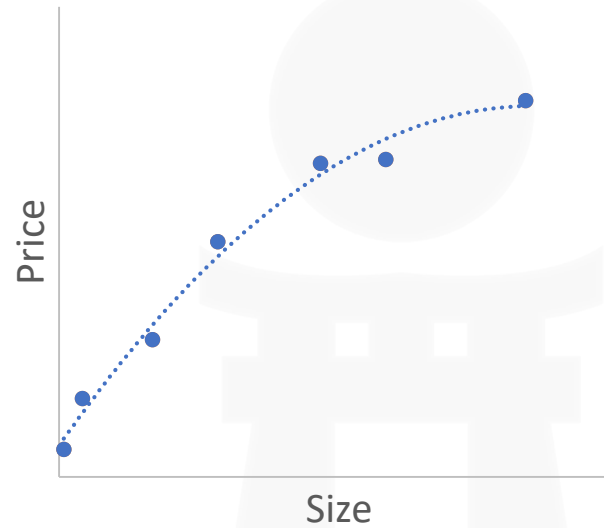
REGULARIZATION




Overfitting



$$\theta_0 + \theta_1 x$$




$$\theta_0 + \theta_1 x + \theta_2 x^2$$




$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

Intuition



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

- Want to discourage complex models automatically – How?
- Adjust the cost function!
 - Penalize models with large high-order θ terms

$$J'(\theta) = J(\theta) + \text{Penalty}$$

Definitions

- Two most common methods

- L1 regularization
 - lasso regression

$$J_{L1}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m |\theta_j|$$

- L2 regularization
 - ridge regression
 - weight decay

$$J_{L2}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m \theta_j^2$$

Regularized Regression

$$J_{L1}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m |\theta_j|$$

$$J_{L2}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m \theta_j^2$$

- Find the best fit
- Keep the θ_j terms as small as possible.
- λ is a user-set parameter which controls the trade off

Regularized Regression

$$J_{L1}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m |\theta_j|$$

$$J_{L2}(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^m \theta_j^2$$

- Size of λ important
 - λ too high \Rightarrow no fitting
 - λ too low \Rightarrow no regularization

QUESTIONS

